



SUPERVISED LEARNING WITH SCIKIT-LEARN

Introduction to regression



Boston housing data

```
In [1]: boston = pd.read_csv('boston.csv')
```

```
In [2]: print(boston.head())
```

	CRIM	ZN	INDUS	CHAS	NX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	

	PTRATIO	B	LSTAT	MEDV
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2



Creating feature and target arrays

```
In [3]: X = boston.drop('MEDV', axis=1).values
```

```
In [4]: y = boston['MEDV'].values
```



Predicting house value from a single feature

```
In [5]: X_rooms = X[:,5]
```

```
In [6]: type(X_rooms), type(y)
```

```
Out[6]: (numpy.ndarray, numpy.ndarray)
```

```
In [7]: y = y.reshape(-1, 1)
```

```
In [8]: X_rooms = X_rooms.reshape(-1, 1)
```



Plotting house value vs. number of rooms

```
In [9]: plt.scatter(X_rooms, y)
```

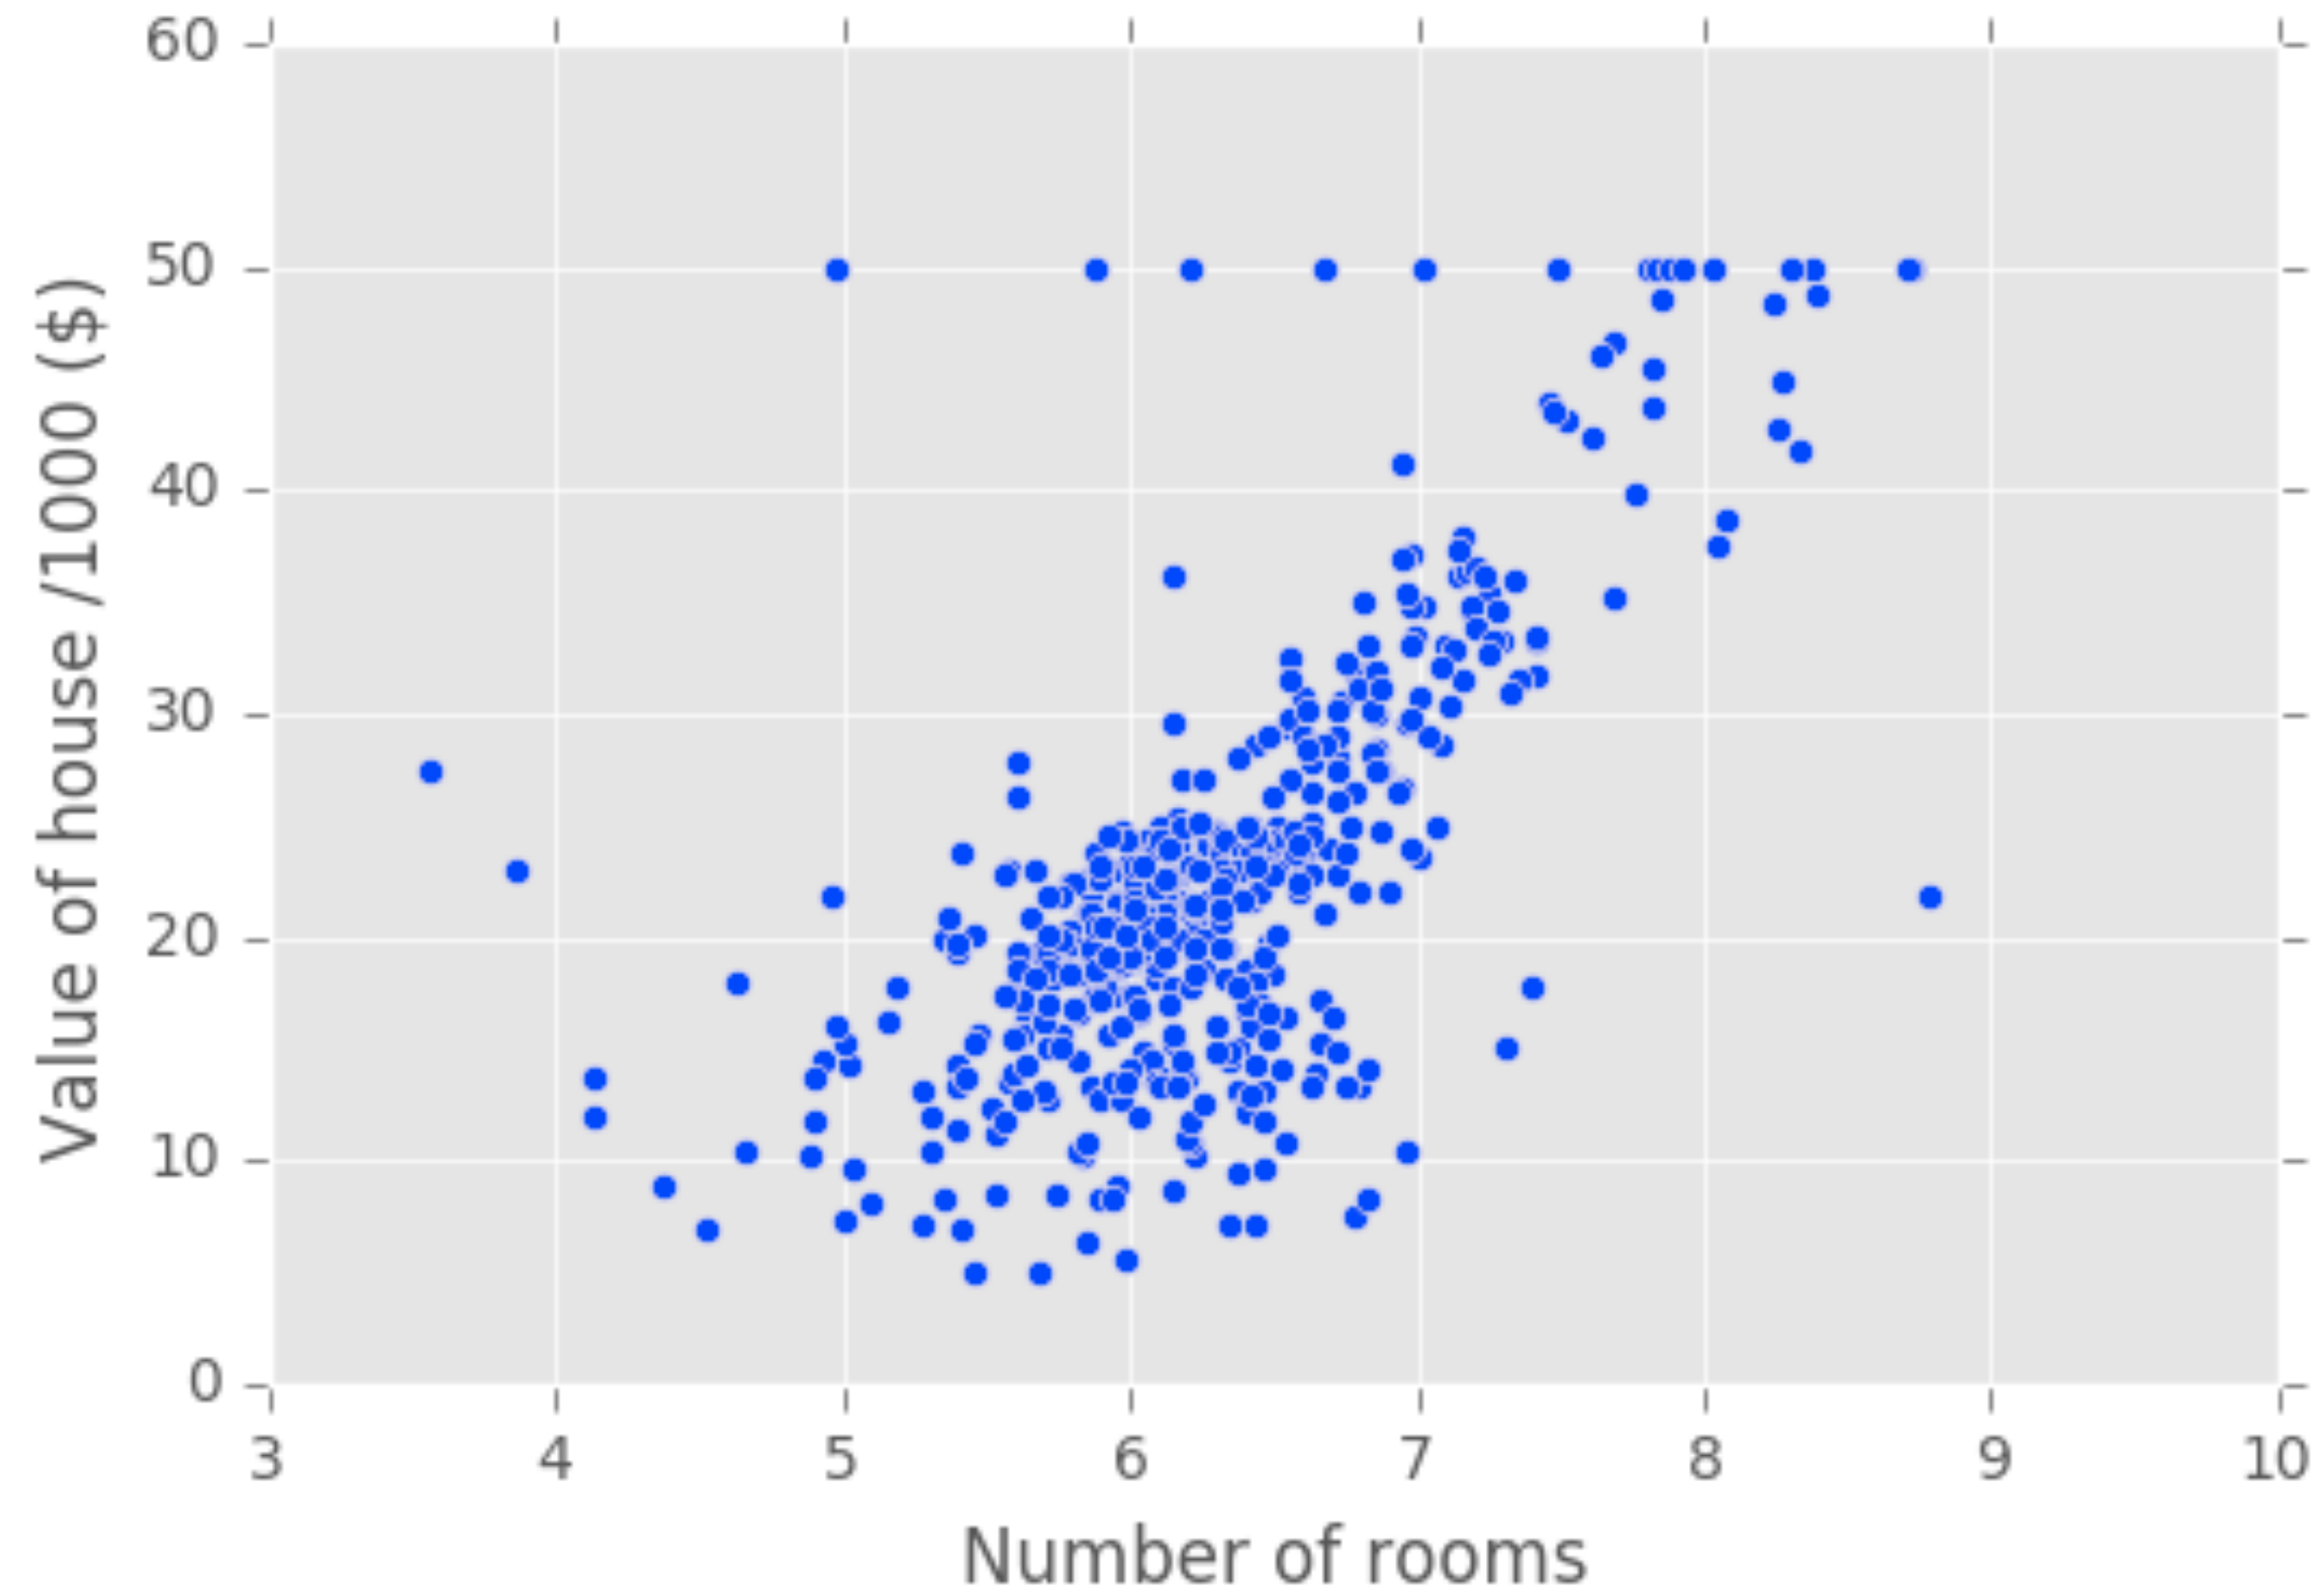
```
In [10]: plt.ylabel('Value of house /1000 ($)')
```

```
In [11]: plt.xlabel('Number of rooms')
```

```
In [12]: plt.show();
```



Plotting house value vs. number of rooms





Fitting a regression model

```
In [13]: import numpy as np
```

```
In [14]: from sklearn import linear_model
```

```
In [15]: reg = linear_model.LinearRegression()
```

```
In [16]: reg.fit(X_rooms, y)
```

```
In [17]: prediction_space = np.linspace(min(X_rooms),  
....:                                   max(X_rooms)).reshape(-1, 1)
```

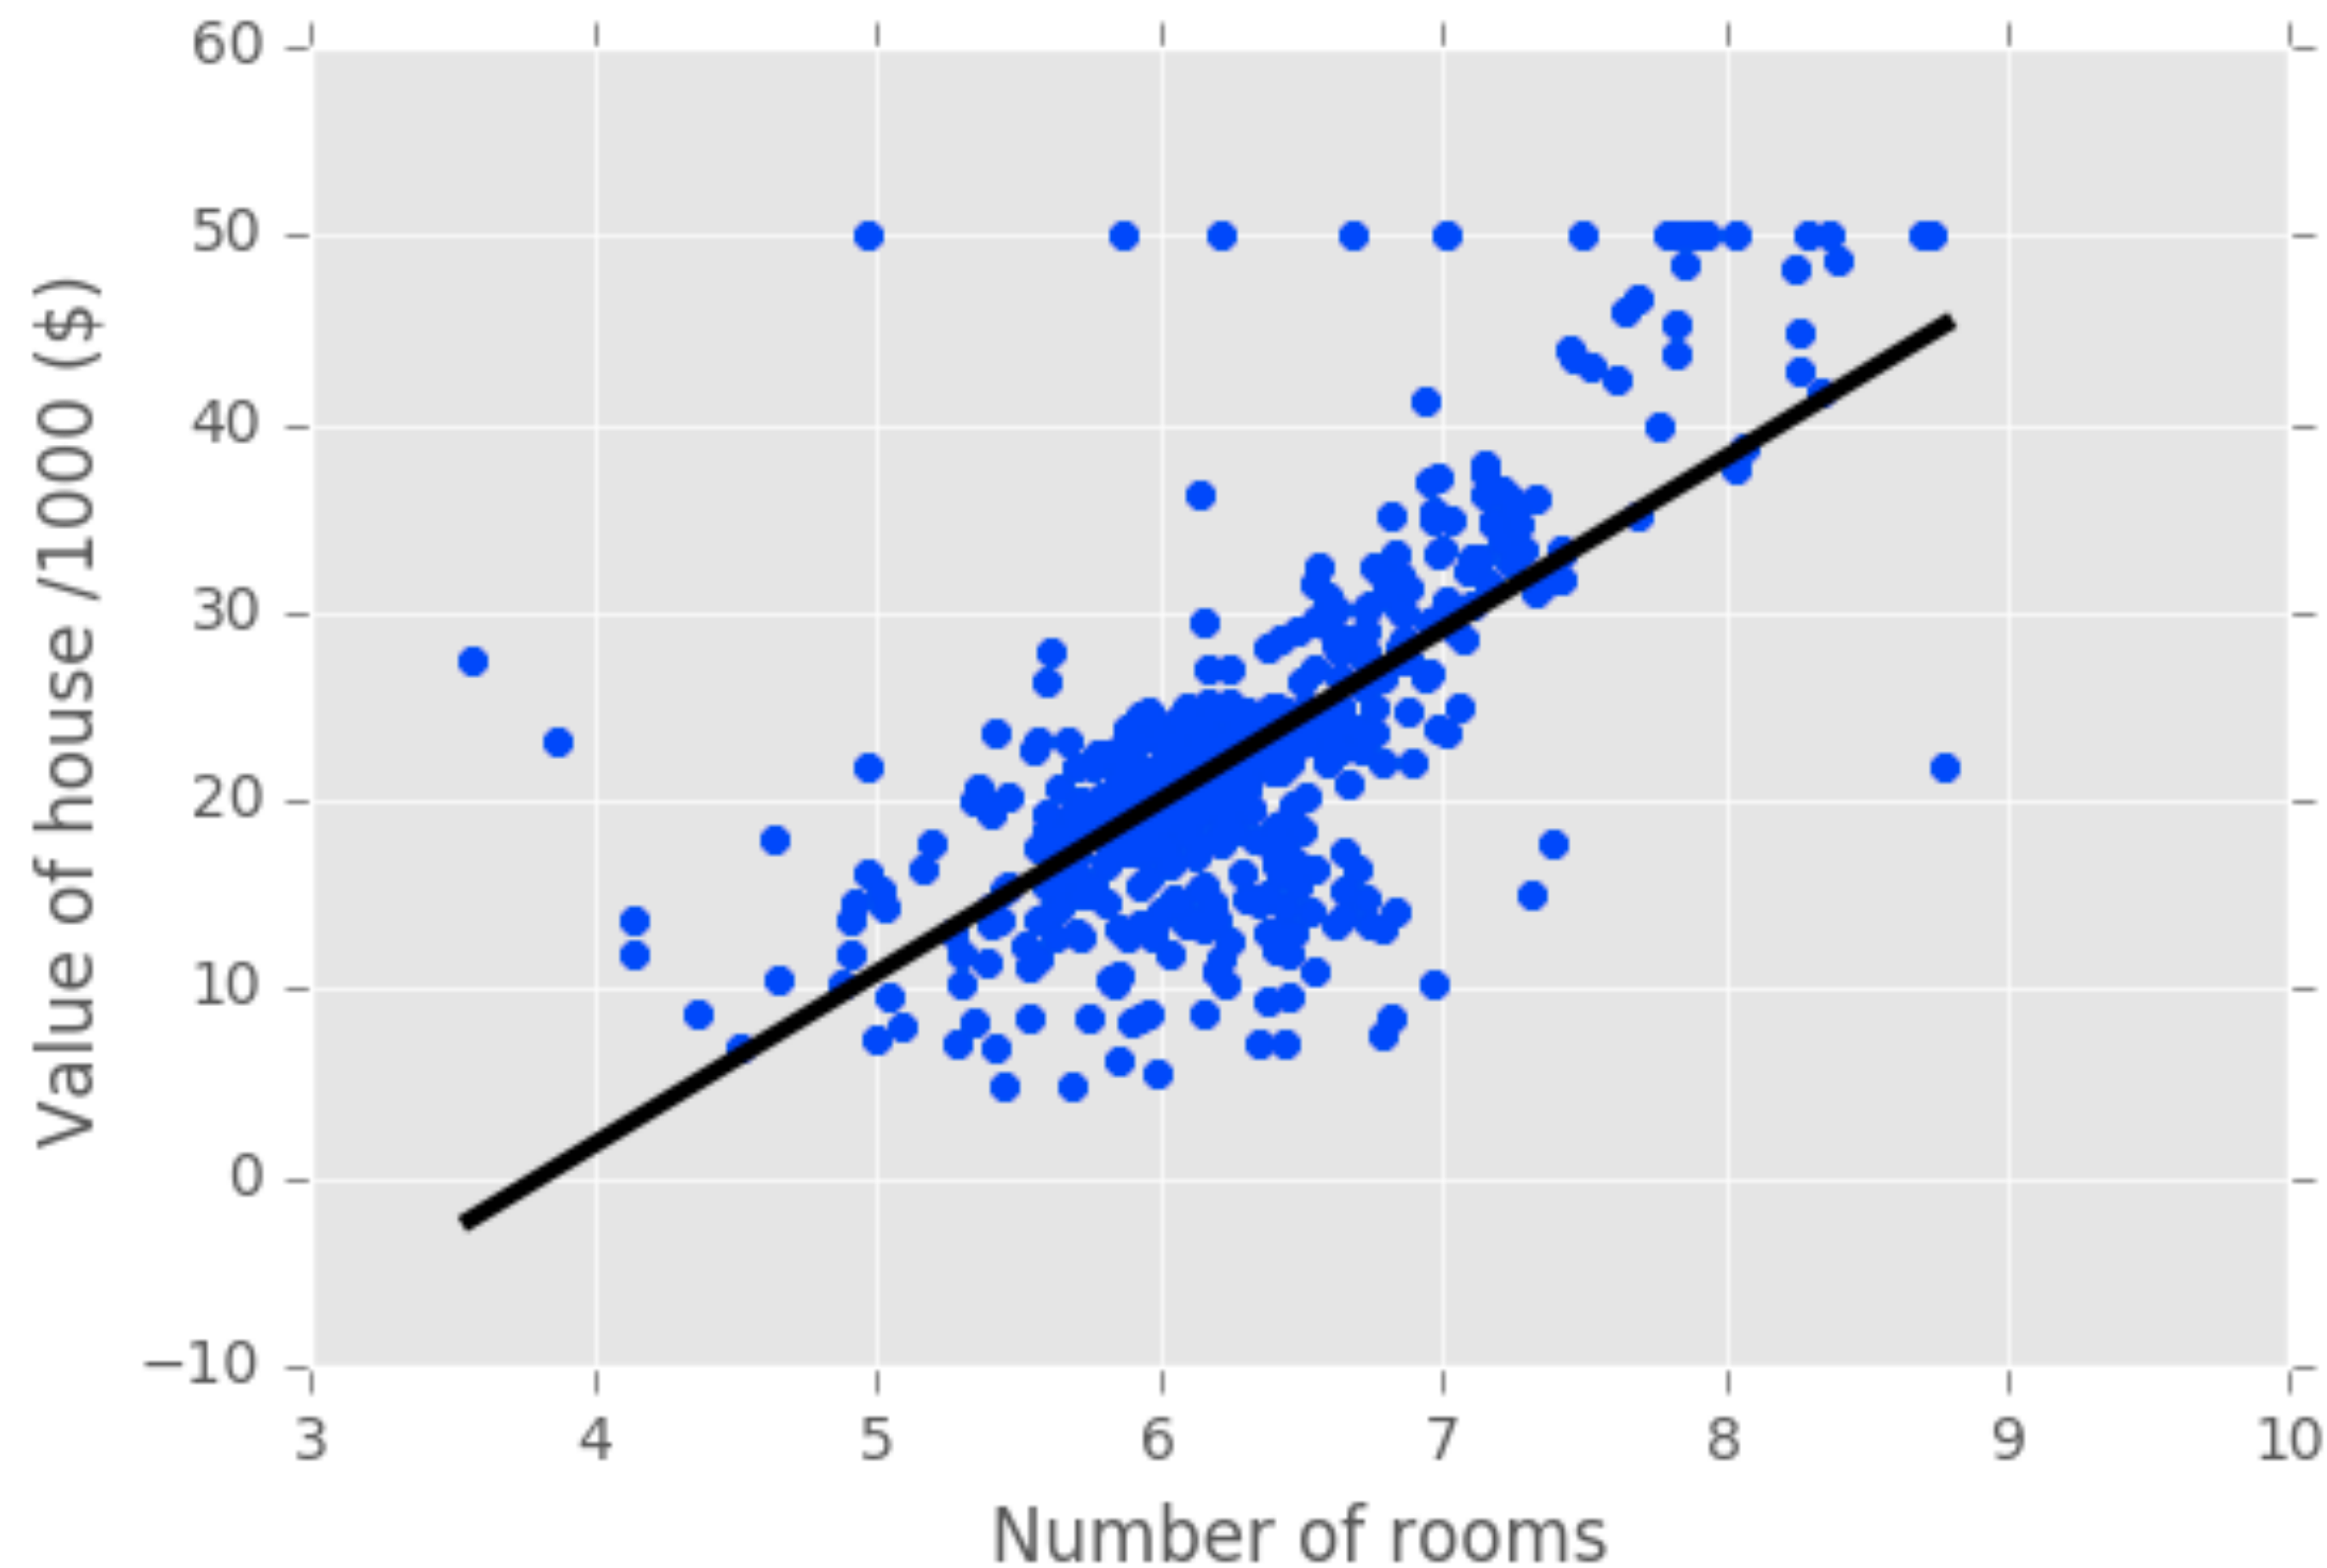
```
In [18]: plt.scatter(X_rooms, y, color='blue')
```

```
In [19]: plt.plot(prediction_space, reg.predict(prediction_space),  
....:              color='black', linewidth=3)
```

```
In [20]: plt.show()
```



Fitting a regression model





SUPERVISED LEARNING WITH SCIKIT-LEARN

Let's practice!



SUPERVISED LEARNING WITH SCIKIT-LEARN

The basics of linear regression

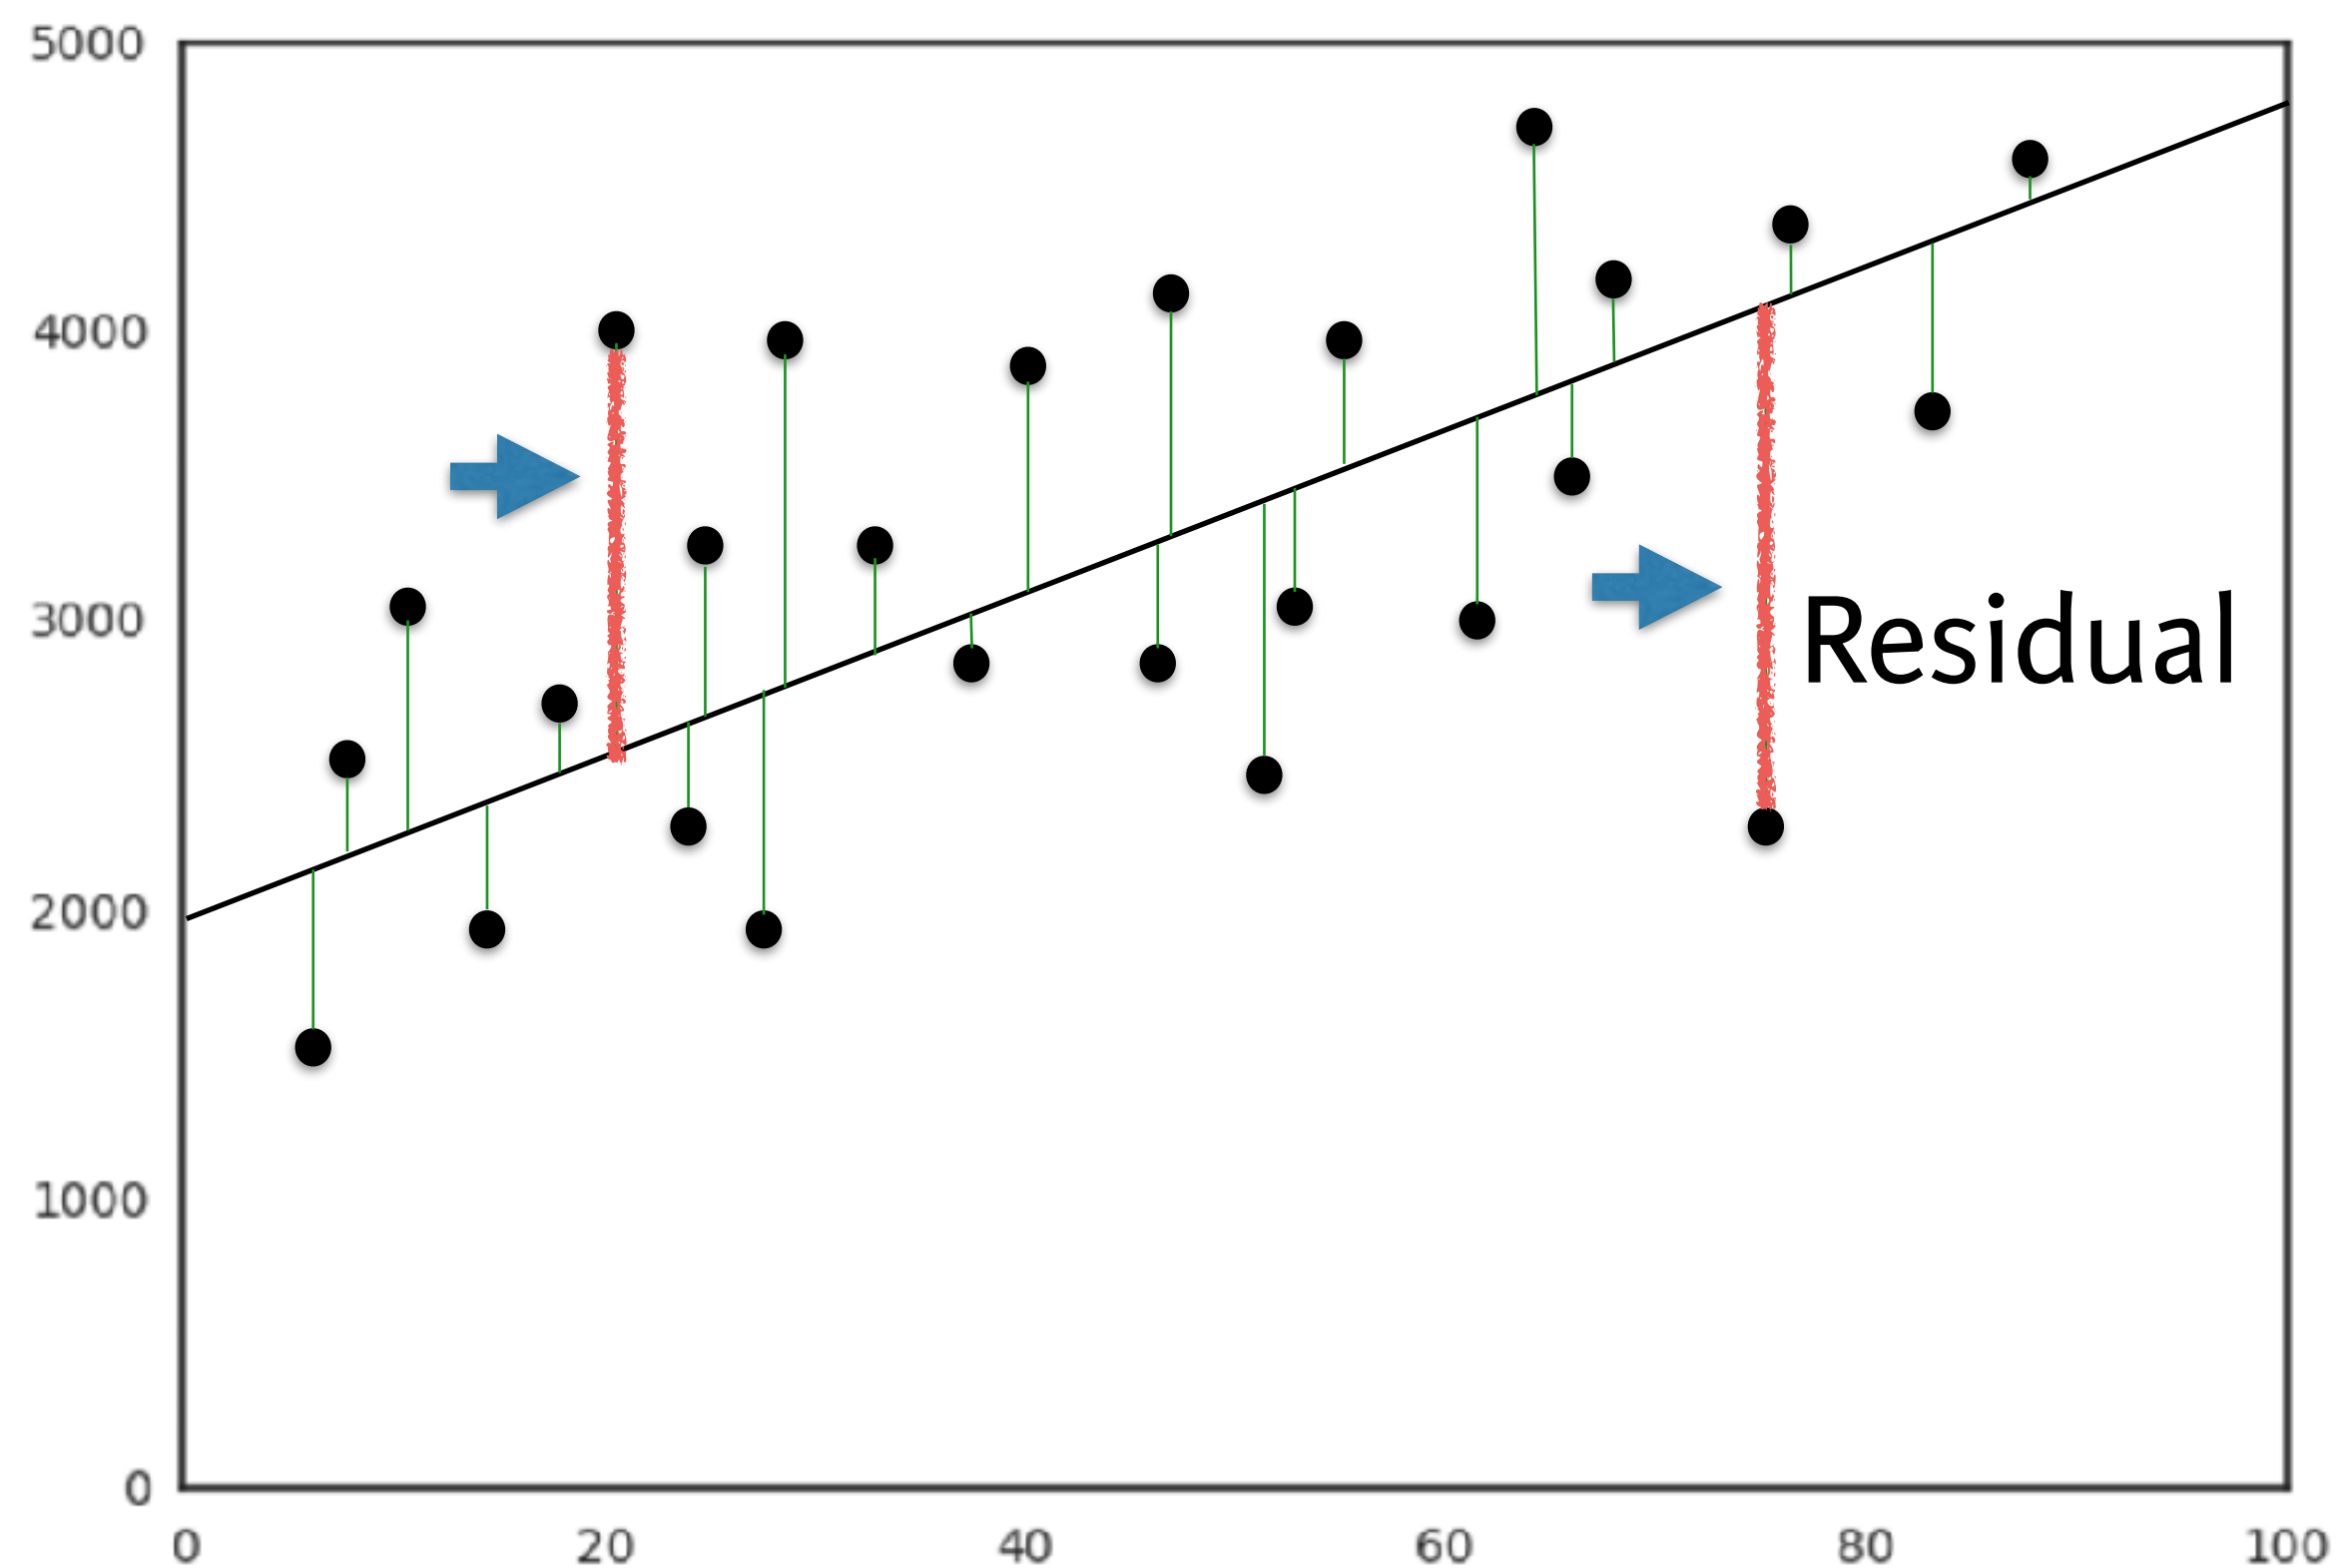


Regression mechanics

- $y = ax + b$
 - $y = \text{target}$
 - $x = \text{single feature}$
 - $a, b = \text{parameters of model}$
- How do we choose a and b ?
- Define an error function for any given line
 - Choose the line that minimizes the error function

The loss function

- Ordinary least squares (OLS): Minimize sum of squares of residuals





Linear regression in higher dimensions

$$y = a_1x_1 + a_2x_2 + b$$

- To fit a linear regression model here:
 - Need to specify 3 variables

- In higher dimensions:

$$y = a_1x_1 + a_2x_2 + a_3x_3 + a_nx_n + b$$

- Must specify coefficient for each feature and the variable b
- Scikit-learn API works exactly the same way:
 - Pass two arrays: Features, and target



Linear regression on all features

```
In [1]: from sklearn.model_selection import train_test_split

In [2]: X_train, X_test, y_train, y_test = train_test_split(X, y,
...: test_size = 0.3, random_state=42)

In [3]: reg_all = linear_model.LinearRegression()

In [4]: reg_all.fit(X_train, y_train)

In [5]: y_pred = reg_all.predict(X_test)

In [6]: reg_all.score(X_test, y_test)
Out[6]: 0.71122600574849526
```



SUPERVISED LEARNING WITH SCIKIT-LEARN

Let's practice!



SUPERVISED LEARNING WITH SCIKIT-LEARN

Cross-validation

Cross-validation motivation

- Model performance is dependent on way the data is split
- Not representative of the model's ability to generalize
- Solution: Cross-validation!



Cross-validation basics

Split 1	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Metric 1
Split 2	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Metric 2
Split 3	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Metric 3
Split 4	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Metric 4
Split 5	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Metric 5

Training data

Test data

Cross-validation and model performance

- 5 folds = 5-fold CV
- 10 folds = 10-fold CV
- k folds = k-fold CV
- More folds = More computationally expensive



Cross-validation in scikit-learn

```
In [1]: from sklearn.model_selection import cross_val_score

In [2]: reg = linear_model.LinearRegression()

In [3]: cv_results = cross_val_score(reg, X, y, cv=5)

In [4]: print(cv_results)
[ 0.63919994  0.71386698  0.58702344  0.07923081 -0.25294154]

In [5]: np.mean(cv_results)
Out[5]: 0.35327592439587058
```



SUPERVISED LEARNING WITH SCIKIT-LEARN

Let's practice!



SUPERVISED LEARNING WITH SCIKIT-LEARN

Regularized regression



Why regularize?

- Recall: Linear regression minimizes a loss function
- It chooses a coefficient for each feature variable
- Large coefficients can lead to overfitting
- Penalizing large coefficients: Regularization



Ridge regression

- Loss function = OLS loss function + $\alpha * \sum_{i=1}^n a_i^2$
- Alpha: Parameter we need to choose
- Picking alpha here is similar to picking k in k-NN
- Hyperparameter tuning (More in Chapter 3)
- Alpha controls model complexity
 - Alpha = 0: We get back OLS (Can lead to overfitting)
 - Very high alpha: Can lead to underfitting



Ridge regression in scikit-learn

```
In [1]: from sklearn.linear_model import Ridge
```

```
In [2]: X_train, X_test, y_train, y_test = train_test_split(X, y,  
...: test_size = 0.3, random_state=42)
```

```
In [3]: ridge = Ridge(alpha=0.1, normalize=True)
```

```
In [4]: ridge.fit(X_train, y_train)
```

```
In [5]: ridge_pred = ridge.predict(X_test)
```

```
In [6]: ridge.score(X_test, y_test)
```

```
Out[6]: 0.69969382751273179
```



Lasso regression

- Loss function = OLS loss function + $\alpha * \sum_{i=1}^n |a_i|$

this function plots the R2 score as well as standard error for each alpha:

```
def display_plot(cv_scores, cv_scores_std):  
    fig = plt.figure()  
    ax = fig.add_subplot(1,1,1)  
    ax.plot(alpha_space, cv_scores)  
    std_error = cv_scores_std / np.sqrt(10)  
    ax.fill_between(alpha_space, cv_scores + std_error, cv_scores - std_error, alpha=0.2)  
    ax.set_ylabel('CV Score +/- Std Error')  
    ax.set_xlabel('Alpha')  
    ax.axhline(np.max(cv_scores), linestyle='--', color='.5')  
    ax.set_xlim([alpha_space[0], alpha_space[-1]])  
    ax.set_xscale('log')  
    plt.show()
```



Lasso regression in scikit-learn

```
In [1]: from sklearn.linear_model import Lasso
```

```
In [2]: X_train, X_test, y_train, y_test = train_test_split(X, y,  
...: test_size = 0.3, random_state=42)
```

```
In [3]: lasso = Lasso(alpha=0.1, normalize=True)
```

```
In [4]: lasso.fit(X_train, y_train)
```

```
In [5]: lasso_pred = lasso.predict(X_test)
```

```
In [6]: lasso.score(X_test, y_test)
```

```
Out[6]: 0.59502295353285506
```



Lasso regression for feature selection

- Can be used to select important features of a dataset
- Shrinks the coefficients of less important features to exactly 0

```
# Import necessary modules
from sklearn.linear_model import Ridge
from sklearn.model_selection import cross_val_score
# Setup the array of alphas and lists to store scores
alpha_space = np.logspace(-4, 0, 50)
ridge_scores = []
ridge_scores_std = []
# Create a ridge regressor:
ridgeridge = Ridge(normalize = True)
# Compute scores over range of alphas
for alpha in alpha_space:
    # Specify the alpha value to use: ridge.alpha
    ridge.alpha = alpha
    # Perform 10-fold CV: ridge_cv_scores
    ridge_cv_scores = cross_val_score(ridge, X, y, cv = 10)
    # Append the mean of ridge_cv_scores to ridge_scores
    ridge_scores.append(np.mean(ridge_cv_scores))
    # Append the std of ridge_cv_scores to ridge_scores_std
    ridge_scores_std.append(np.std(ridge_cv_scores))
# Display the plot

display_plot(ridge_scores, ridge_scores_std)
```



Lasso for feature selection in scikit-learn

```
In [1]: from sklearn.linear_model import Lasso
```

```
In [2]: names = boston.drop('MEDV', axis=1).columns
```

```
In [3]: lasso = Lasso(alpha=0.1)
```

```
In [4]: lasso_coef = lasso.fit(X, y).coef_
```

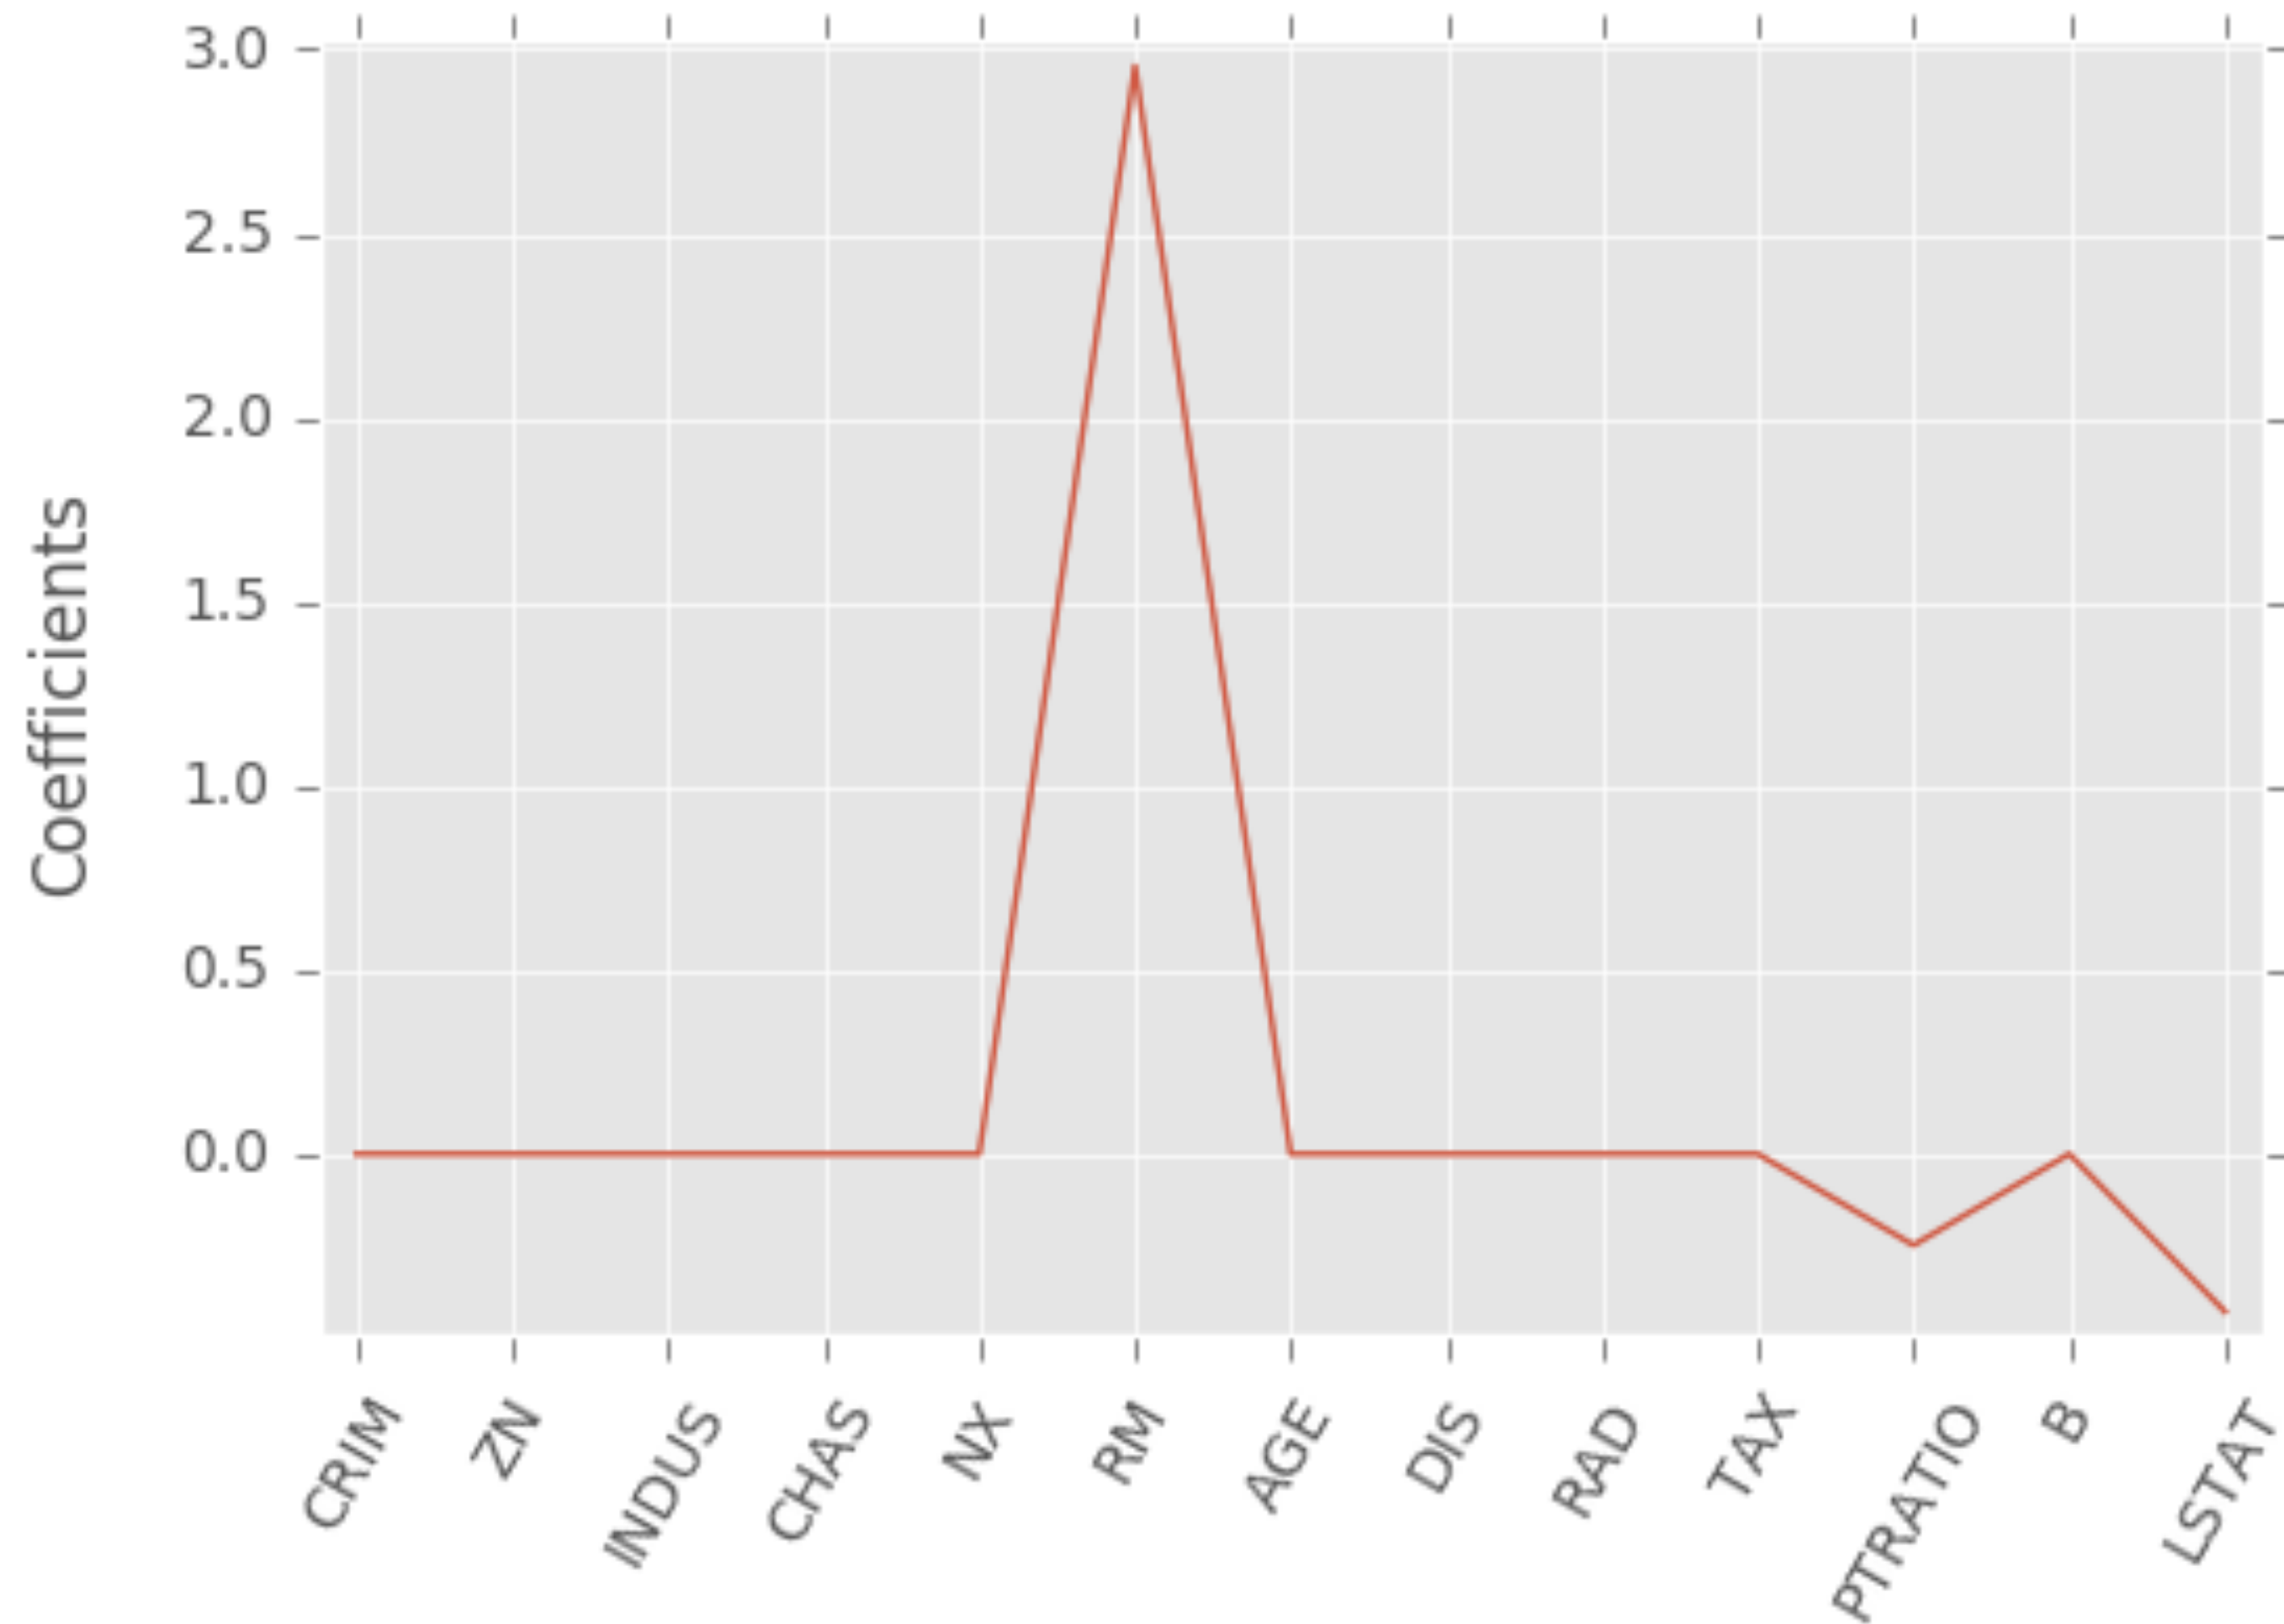
```
In [5]: _ = plt.plot(range(len(names)), lasso_coef)
```

```
In [6]: _ = plt.xticks(range(len(names)), names, rotation=60)
```

```
In [7]: _ = plt.ylabel('Coefficients')
```

```
In [8]: plt.show()
```

Lasso for feature selection in scikit-learn





SUPERVISED LEARNING WITH SCIKIT-LEARN

Let's practice!